

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV, StratifiedKFold
from sklearn.preprocessing import StandardScaler, OneHotEncoder, LabelEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, f1_score
from sklearn.preprocessing import LabelEncoder, StandardScaler, OneHotEncoder # Added OneHotEncoder
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier # Example of another model
from sklearn.metrics import accuracy_score, classification_report
from sklearn.compose import ColumnTransformer # For robust preprocessing
from sklearn.pipeline import Pipeline
import pickle # For saving/loading
import joblib # Often preferred for scikit-learn models

sns.set_style("whitegrid")

# Set style for plots
%matplotlib inline
sns.set_style("whitegrid")

```

▼ Load the datasets

```

try:
    crop_df = pd.read_csv('Crop_recommendation.csv')
    fertilizer_df = pd.read_csv('Fertilizer Prediction.csv')
    print("Datasets loaded successfully!")
except FileNotFoundError:
    print("Make sure 'Crop_recommendation.csv' and 'Fertilizer Prediction.csv' are in the same directory.")
    # As a fallback, let's use the content provided in the prompt if files are not found
    crop_data_string = """N,P,K,temperature,humidity,ph,rainfall,label
90,42,43,20.87974371,82.00274423,6.502985292000001,202.9355362,rice
... (rest of the crop data) ...
120,16,51,27.99901833,91.64193051,6.547041902999999,23.28618248,muskmelon""" # Truncated for brevity

    fertilizer_data_string = """Temparature,Humidity ,Moisture,Soil Type,Crop Type,Nitrogen,Potassium,Phosphorous,Fertilizer Name
26,52,38,Sandy,Maize,37,0,0,Urea
... (rest of the fertilizer data) ...
29,58,57,Black,Sugarcane,12,0,10,20-20""" # Truncated for brevity

    from io import StringIO
    if 'crop_df' not in locals():
        try:
            crop_df = pd.read_csv(StringIO(crop_data_string)) # Replace with full data string
            print("Crop data loaded from string.")
        except Exception as e:
            print(f"Error loading crop data from string: {e}")
            print("Please ensure the full CSV content is pasted if using string loading.")

    if 'fertilizer_df' not in locals():
        try:
            fertilizer_df = pd.read_csv(StringIO(fertilizer_data_string)) # Replace with full data string
            print("Fertilizer data loaded from string.")
            # Clean up column names if needed
            fertilizer_df.columns = fertilizer_df.columns.str.strip()
        except Exception as e:
            print(f"Error loading fertilizer data from string: {e}")
            print("Please ensure the full CSV content is pasted if using string loading.")

# If files were loaded, fertilizer_df might have leading/trailing spaces in column names
if 'fertilizer_df' in locals() and hasattr(fertilizer_df, 'columns'):
    fertilizer_df.columns = fertilizer_df.columns.str.strip()

```

➡ Datasets loaded successfully!

Project 1: Crop Recommendation System

Goal: Predict the label (crop type) given soil nutrient levels (N, P, K) and environmental conditions (temperature, humidity, pH, rainfall). This is a multi-class classification problem.

✧ Exploratory Data Analysis (EDA) for Crop Recommendation

```
print("\n--- Crop Recommendation Dataset ---")
if 'crop_df' in locals():
    print("First 5 rows:\n", crop_df.head())
    print("\nShape of data:", crop_df.shape)
    print("\nData types:\n", crop_df.dtypes) # All seem numeric except label
    print("\nMissing values:\n", crop_df.isnull().sum()) # Usually no missing values in this dataset

    print("\nDescriptive statistics:\n", crop_df.describe())

    print("\nCrop Labels Distribution:")
    print(crop_df['label'].value_counts())
    plt.figure(figsize=(12, 6))
    sns.countplot(y='label', data=crop_df, order=crop_df['label'].value_counts().index)
    plt.title('Distribution of Crop Labels')
    plt.tight_layout()
    plt.show()

    # Histograms for numerical features
    crop_df.drop('label', axis=1).hist(figsize=(12, 10), bins=20)
    plt.suptitle('Histograms of Numerical Features (Crop Data)', y=1.02)
    plt.tight_layout()
    plt.show()

    # Boxplots to see feature distribution per crop
    numerical_cols_crop = crop_df.select_dtypes(include=np.number).columns
    for col in numerical_cols_crop:
        plt.figure(figsize=(15, 7))
        sns.boxplot(x='label', y=col, data=crop_df)
        plt.title(f'{col} by Crop Type')
        plt.xticks(rotation=90)
        plt.tight_layout()
        plt.show()

    # Correlation heatmap
    plt.figure(figsize=(10, 8))
    sns.heatmap(crop_df[numerical_cols_crop].corr(), annot=True, cmap='coolwarm', fmt=".2f")
    plt.title('Correlation Matrix (Crop Data)')
    plt.show()
else:
    print("crop_df not loaded. Skipping EDA for Crop Recommendation.")
```



--- Crop Recommendation Dataset ---

First 5 rows:

	N	P	K	temperature	humidity	ph	rainfall	label
0	90	42	43	20.879744	82.002744	6.502985	202.935536	rice
1	85	58	41	21.770462	80.319644	7.038096	226.655537	rice
2	60	55	44	23.004459	82.320763	7.840207	263.964248	rice
3	74	35	40	26.491096	80.158363	6.980401	242.864034	rice
4	78	42	42	20.130175	81.604873	7.628473	262.717340	rice

Shape of data: (2200, 8)

Data types:

N	int64
P	int64
K	int64
temperature	float64
humidity	float64
ph	float64
rainfall	float64
label	object

dtype: object

Missing values:

N	0
P	0
K	0
temperature	0
humidity	0
ph	0
rainfall	0
label	0

dtype: int64

Descriptive statistics:

	N	P	K	temperature	humidity \
count	2200.000000	2200.000000	2200.000000	2200.000000	2200.000000
mean	50.551818	53.362727	48.149091	25.616244	71.481779
std	36.917334	32.985883	50.647931	5.063749	22.263812
min	0.000000	5.000000	5.000000	8.825675	14.258040
25%	21.000000	28.000000	20.000000	22.769375	60.261953
50%	37.000000	51.000000	32.000000	25.598693	80.473146
75%	84.250000	68.000000	49.000000	28.561654	89.948771
max	140.000000	145.000000	205.000000	43.675493	99.981876

	ph	rainfall
count	2200.000000	2200.000000
mean	6.469480	103.463655
std	0.773938	54.958389
min	3.504752	20.211267
25%	5.971693	64.551686
50%	6.425045	94.867624
75%	6.923643	124.267508
max	9.935091	298.560117


Crop Labels Distribution:

label	
rice	100
maize	100
chickpea	100
kidneybeans	100
pigeonpeas	100
mothbeans	100
mungbean	100
blackgram	100
lentil	100
pomegranate	100
banana	100
mango	100
grapes	100
watermelon	100
muskmelon	100
apple	100
orange	100
papaya	100
coconut	100
cotton	100
jute	100
coffee	100

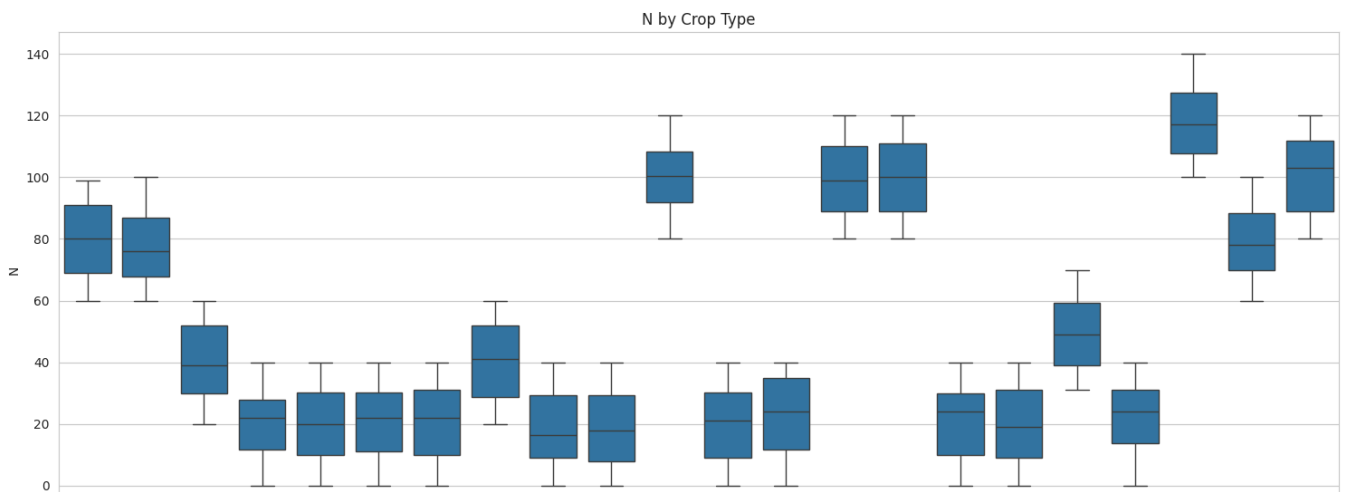
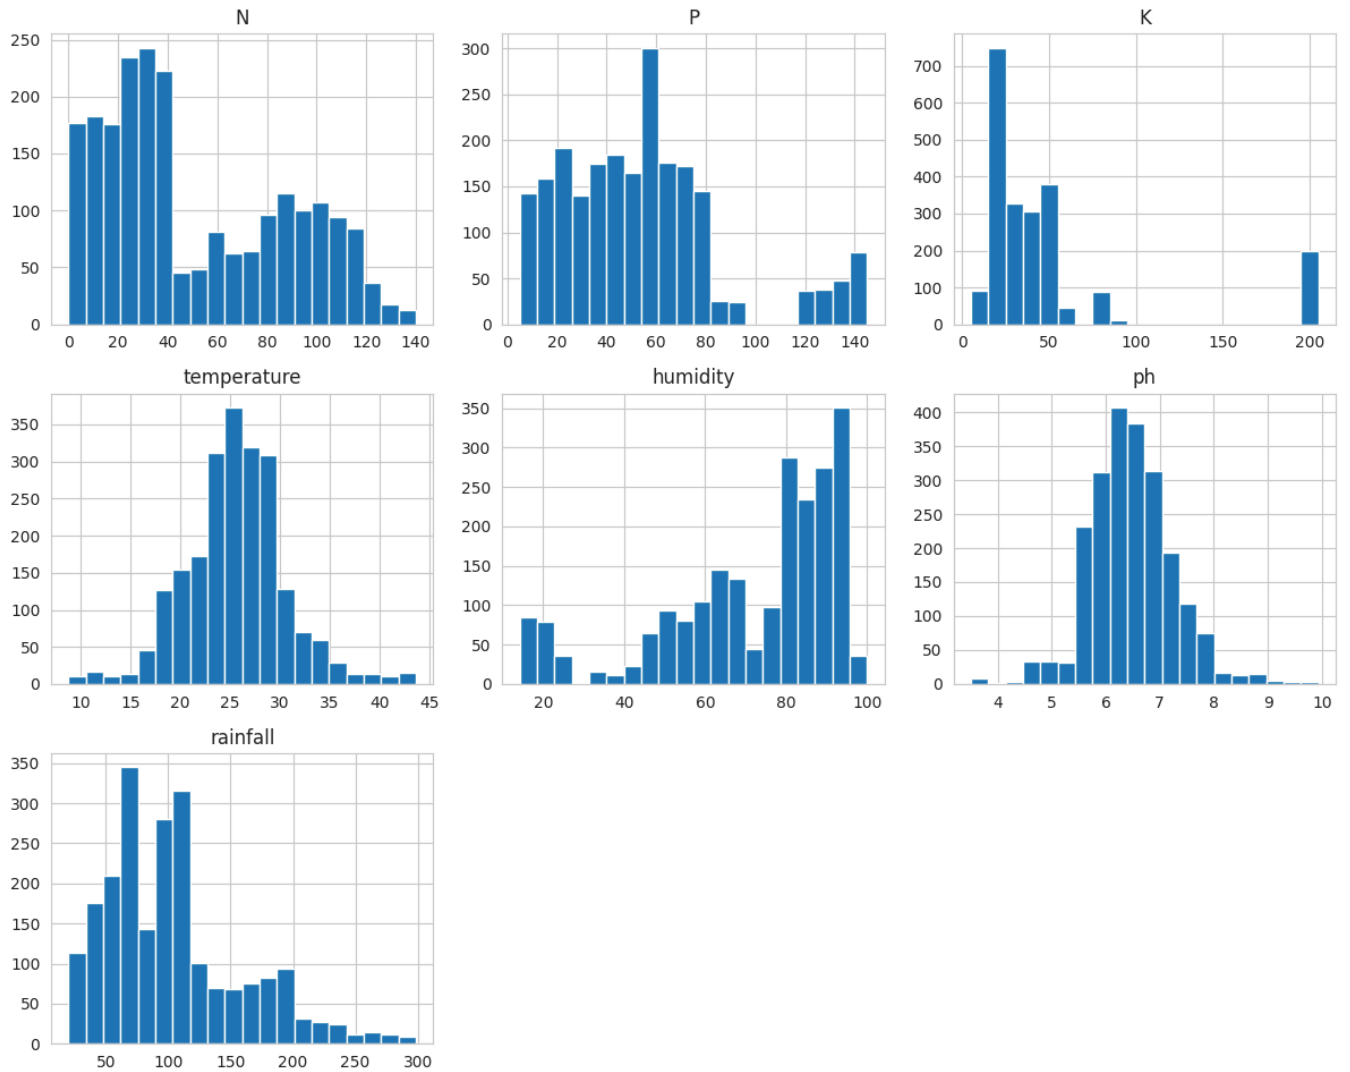
Name: count, dtype: int64


Distribution of Crop Labels


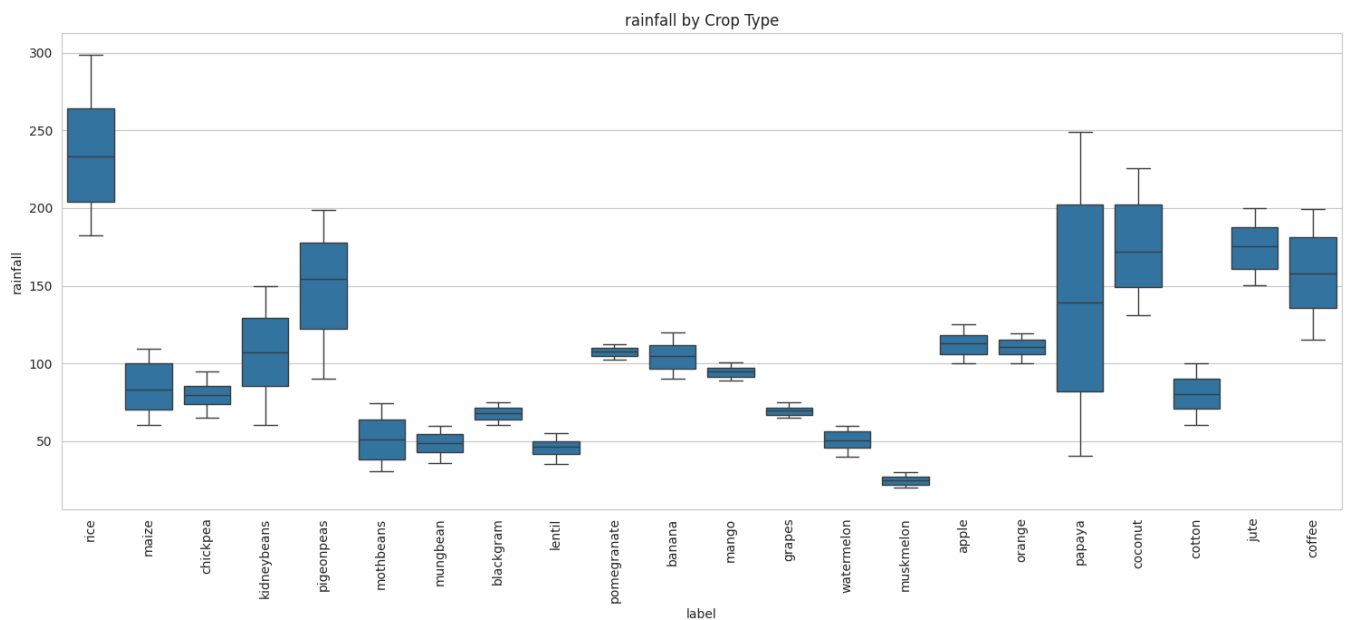
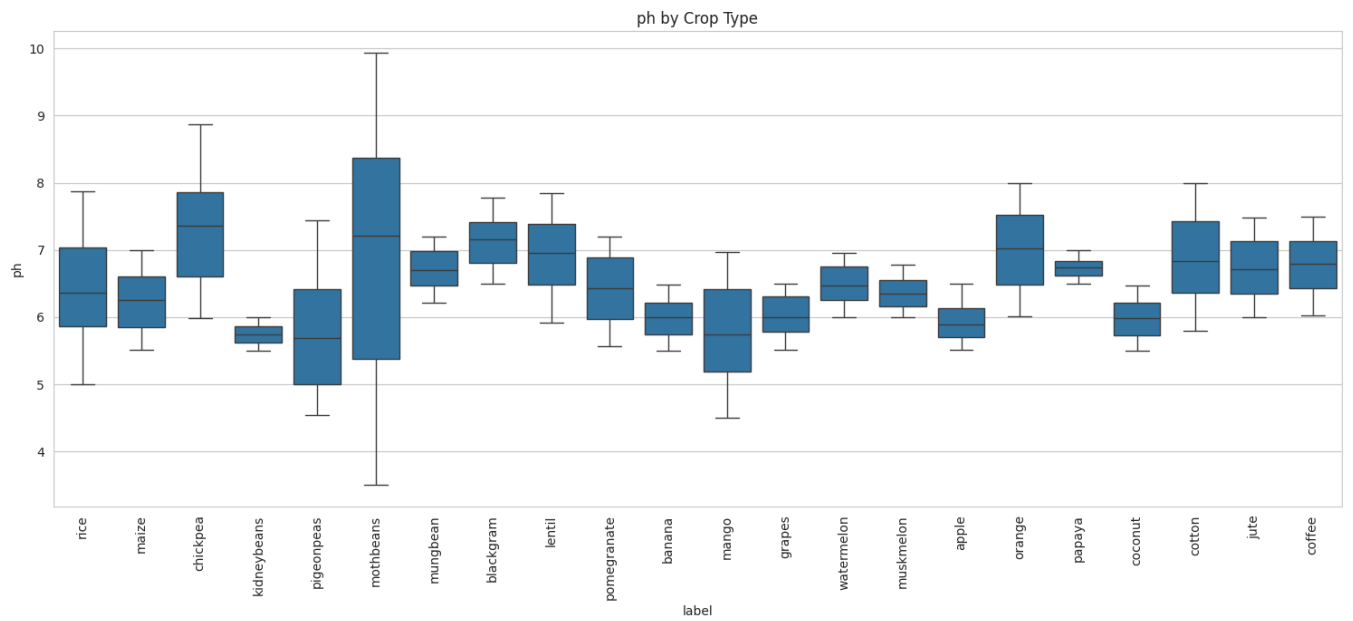



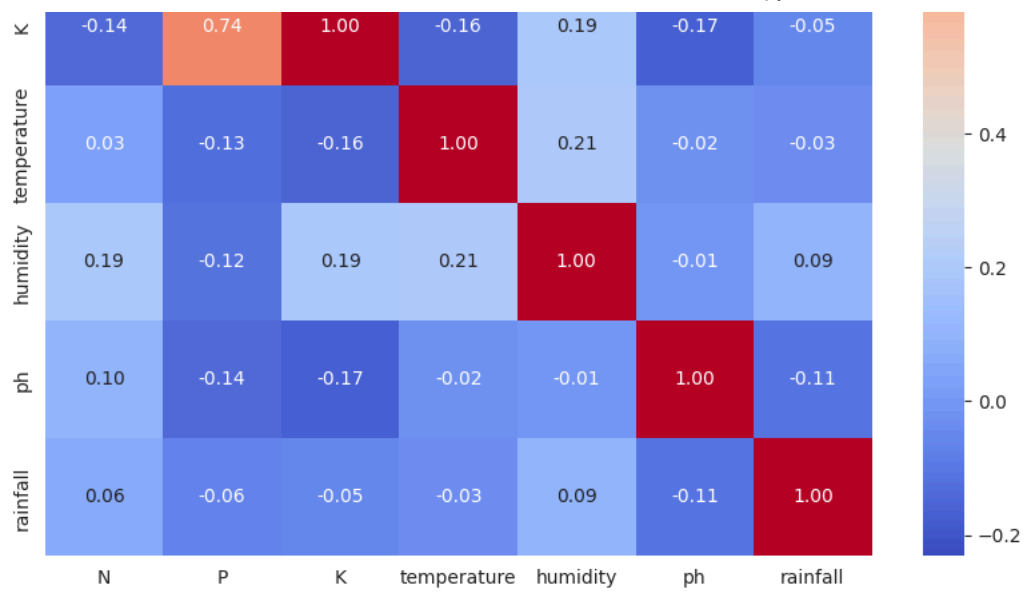


Histograms of Numerical Features (Crop Data)









2. Data Preprocessing for Crop Recommendation

```

if 'crop_df' in locals():
    X_crop = crop_df.drop('label', axis=1)
    y_crop_str = crop_df['label']

    # Encode the target variable (crop labels)
    label_encoder_crop = LabelEncoder()
    y_crop = label_encoder_crop.fit_transform(y_crop_str)
    print("\nEncoded Crop Labels Mapping:")
    for i, class_name in enumerate(label_encoder_crop.classes_):
        print(f"{class_name}: {i}")

    # Split data into training and testing sets
    X_train_crop, X_test_crop, y_train_crop, y_test_crop = train_test_split(
        X_crop, y_crop, test_size=0.2, random_state=42, stratify=y_crop
    )

    # Feature Scaling (Numerical features)
    # All features are numerical, so we can apply StandardScaler to all of X
    scaler_crop = StandardScaler()
    X_train_crop_scaled = scaler_crop.fit_transform(X_train_crop)
    X_test_crop_scaled = scaler_crop.transform(X_test_crop)

    print("\nShape of X_train_crop_scaled:", X_train_crop_scaled.shape)
    print("Shape of X_test_crop_scaled:", X_test_crop_scaled.shape)
else:
    print("crop_df not loaded. Skipping Preprocessing for Crop Recommendation.")

```



Encoded Crop Labels Mapping:

```

apple: 0
banana: 1
blackgram: 2
chickpea: 3
coconut: 4
coffee: 5
cotton: 6
grapes: 7
jute: 8
kidneybeans: 9
lentil: 10
maize: 11
mango: 12
mothbeans: 13
mungbean: 14
muskmelon: 15
orange: 16
papaya: 17
pigeonpeas: 18
pomegranate: 19
rice: 20
watermelon: 21

```

```

Shape of X_train_crop_scaled: (1760, 7)
Shape of X_test_crop_scaled: (440, 7)

```

3. Model Selection & Training for Crop Recommendation

```

if 'crop_df' in locals():
    models_crop = {
        "Logistic Regression": LogisticRegression(max_iter=1000, solver='liblinear', multi_class='ovr', random_state=42),
        "K-Nearest Neighbors": KNeighborsClassifier(),
        "Decision Tree": DecisionTreeClassifier(random_state=42),
        "Random Forest": RandomForestClassifier(random_state=42),
        "Gradient Boosting": GradientBoostingClassifier(random_state=42)
    }

    results_crop = {}
    print("\n--- Training Crop Recommendation Models ---")

    for name, model in models_crop.items():
        model.fit(X_train_crop_scaled, y_train_crop)
        y_pred_crop = model.predict(X_test_crop_scaled)

        accuracy = accuracy_score(y_test_crop, y_pred_crop)
        # For multi-class, 'weighted' F1 is often a good summary
        f1 = f1_score(y_test_crop, y_pred_crop, average='weighted')
        results_crop[name] = {'Accuracy': accuracy, 'F1-score (Weighted)': f1}

    print(f"\n--- {name} (Crop Recommendation) ---")

```



```

print(f"Accuracy: {accuracy:.4f}")
print(f"F1-score (Weighted): {f1:.4f}")
# print("Classification Report:\n", classification_report(y_test_crop, y_pred_crop, target_names=label_encoder_crop.classes_, ze
# Detailed report can be long, so commented out for brevity here.

crop_results_df = pd.DataFrame(results_crop).T.sort_values(by='Accuracy', ascending=False)
print("\n--- Model Performance Summary (Crop Recommendation) ---")
print(crop_results_df)
else:
    print("crop_df not loaded. Skipping Model Training for Crop Recommendation.")

```

--- Training Crop Recommendation Models ---

--- Logistic Regression (Crop Recommendation) ---
Accuracy: 0.9386
F1-score (Weighted): 0.9382

--- K-Nearest Neighbors (Crop Recommendation) ---
Accuracy: 0.9795
F1-score (Weighted): 0.9793

--- Decision Tree (Crop Recommendation) ---
Accuracy: 0.9795
F1-score (Weighted): 0.9794
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:1256: FutureWarning: 'multi_class' was deprecated in versi
warnings.warn()

--- Random Forest (Crop Recommendation) ---
Accuracy: 0.9955
F1-score (Weighted): 0.9955

--- Gradient Boosting (Crop Recommendation) ---
Accuracy: 0.9886
F1-score (Weighted): 0.9887

--- Model Performance Summary (Crop Recommendation) ---

	Accuracy	F1-score (Weighted)
Random Forest	0.995455	0.995452
Gradient Boosting	0.988636	0.988723
Decision Tree	0.979545	0.979423
K-Nearest Neighbors	0.979545	0.979283
Logistic Regression	0.938636	0.938216

4. Evaluation for Crop Recommendation

```

if 'crop_df' in locals():
    # Let's pick the best model based on accuracy (e.g., Random Forest or Gradient Boosting usually do well)
    # For demonstration, let's assume Random Forest is the best from the results_crop
    best_model_name_crop = crop_results_df.index[0]
    best_model_crop = models_crop[best_model_name_crop] # The trained model instance

    # Re-fit if not already (or get from a dictionary of trained models)
    # The models in `models_crop` are already fitted.
    # If you were to retrain the best one:
    # best_model_crop.fit(X_train_crop_scaled, y_train_crop)

    y_pred_best_crop = best_model_crop.predict(X_test_crop_scaled)

    print(f"\n--- Detailed Evaluation for {best_model_name_crop} (Crop Recommendation) ---")
    print("Classification Report:\n", classification_report(y_test_crop, y_pred_best_crop, target_names=label_encoder_crop.classes_, ze

    cm_crop = confusion_matrix(y_test_crop, y_pred_best_crop)
    plt.figure(figsize=(12, 10))
    sns.heatmap(cm_crop, annot=True, fmt='d', cmap='Blues',
                xticklabels=label_encoder_crop.classes_,
                yticklabels=label_encoder_crop.classes_)
    plt.title(f'Confusion Matrix for {best_model_name_crop} (Crop Recommendation)')
    plt.xlabel('Predicted Label')
    plt.ylabel('True Label')
    plt.xticks(rotation=45, ha='right')
    plt.yticks(rotation=0)
    plt.tight_layout()
    plt.show()
else:
    print("crop_df not loaded. Skipping Evaluation for Crop Recommendation.")

```




--- Detailed Evaluation for Random Forest (Crop Recommendation) ---

Classification Report:

	precision	recall	f1-score	support
apple	1.00	1.00	1.00	20
banana	1.00	1.00	1.00	20
blackgram	1.00	0.95	0.97	20
chickpea	1.00	1.00	1.00	20
coconut	1.00	1.00	1.00	20
coffee	1.00	1.00	1.00	20
cotton	1.00	1.00	1.00	20
grapes	1.00	1.00	1.00	20
jute	0.95	1.00	0.98	20
kidneybeans	1.00	1.00	1.00	20
lentil	1.00	1.00	1.00	20
maize	0.95	1.00	0.98	20
mango	1.00	1.00	1.00	20
mothbeans	1.00	1.00	1.00	20
mungbean	1.00	1.00	1.00	20
muskmelon	1.00	1.00	1.00	20
orange	1.00	1.00	1.00	20
papaya	1.00	1.00	1.00	20
pigeonpeas	1.00	1.00	1.00	20
pomegranate	1.00	1.00	1.00	20
rice	1.00	0.95	0.97	20
watermelon	1.00	1.00	1.00	20
accuracy			1.00	440
macro avg	1.00	1.00	1.00	440
weighted avg	1.00	1.00	1.00	440

Confusion Matrix for Random Forest (Crop Recommendation)



✓ Project 2: Fertilizer Recommendation System

1. Exploratory Data Analysis (EDA) for Fertilizer Recommendation

```

print("\n\n--- Fertilizer Recommendation Dataset ---")
if 'fertilizer_df' in locals():
    print("First 5 rows:\n", fertilizer_df.head())
    print("\nShape of data:", fertilizer_df.shape)
    print("\nData types:\n", fertilizer_df.dtypes)
    print("\nMissing values:\n", fertilizer_df.isnull().sum()) # Usually no missing values

    print("\nDescriptive statistics (numerical features):\n", fertilizer_df.describe())
    print("\nDescriptive statistics (categorical features):\n", fertilizer_df.describe(include='object'))

    print("\nFertilizer Name Distribution:")
    print(fertilizer_df['Fertilizer Name'].value_counts())
    plt.figure(figsize=(10, 5))
    sns.countplot(y='Fertilizer Name', data=fertilizer_df, order=fertilizer_df['Fertilizer Name'].value_counts().index)
    plt.title('Distribution of Fertilizer Names')
    plt.tight_layout()
    plt.show()

    print("\nSoil Type Distribution:")
    print(fertilizer_df['Soil Type'].value_counts())
    sns.countplot(x='Soil Type', data=fertilizer_df)
    plt.title('Distribution of Soil Types')
    plt.show()

    print("\nCrop Type Distribution:")
    print(fertilizer_df['Crop Type'].value_counts())
    plt.figure(figsize=(10,5))
    sns.countplot(y='Crop Type', data=fertilizer_df, order=fertilizer_df['Crop Type'].value_counts().index)
    plt.title('Distribution of Crop Types (Fertilizer Data)')
    plt.tight_layout()
    plt.show()

    # Histograms for numerical features
    numerical_cols_fert = fertilizer_df.select_dtypes(include=np.number).columns
    fertilizer_df[numerical_cols_fert].hist(figsize=(12, 10), bins=15)
    plt.suptitle('Histograms of Numerical Features (Fertilizer Data)', y=1.02)
    plt.tight_layout()
    plt.show()

    # Boxplots for numerical features vs Fertilizer Name
    for col in numerical_cols_fert:
        plt.figure(figsize=(12, 6))
        sns.boxplot(x='Fertilizer Name', y=col, data=fertilizer_df)
        plt.title(f'{col} by Fertilizer Name')
        plt.xticks(rotation=45, ha='right')
        plt.tight_layout()
        plt.show()

    # Categorical features vs Fertilizer Name
    for cat_col in ['Soil Type', 'Crop Type']:
        plt.figure(figsize=(12, 6))
        sns.countplot(x=cat_col, hue='Fertilizer Name', data=fertilizer_df)
        plt.title(f'{cat_col} vs Fertilizer Name')
        plt.xticks(rotation=45, ha='right')
        plt.legend(title='Fertilizer Name', bbox_to_anchor=(1.05, 1), loc='upper left')
        plt.tight_layout()
        plt.show()
else:
    print("fertilizer_df not loaded. Skipping EDA for Fertilizer Recommendation.")

```



--- Fertilizer Recommendation Dataset ---

First 5 rows:

	Temperature	Humidity	Moisture	Soil Type	Crop Type	Nitrogen	Potassium	\
0	26	52	38	Sandy	Maize	37	0	
1	29	52	45	Loamy	Sugarcane	12	0	
2	34	65	62	Black	Cotton	7	9	
3	32	62	34	Red	Tobacco	22	0	
4	28	54	46	Clayey	Paddy	35	0	

	Phosphorous	Fertilizer Name
0	0	Urea
1	36	DAP
2	30	14-35-14
3	20	28-28
4	0	Urea

Shape of data: (99, 9)

Data types:

Temperature	int64
Humidity	int64
Moisture	int64
Soil Type	object
Crop Type	object
Nitrogen	int64
Potassium	int64
Phosphorous	int64
Fertilizer Name	object

dtype: object

Missing values:

Temperature	0
Humidity	0
Moisture	0
Soil Type	0
Crop Type	0
Nitrogen	0
Potassium	0
Phosphorous	0
Fertilizer Name	0

dtype: int64

Descriptive statistics (numerical features):

	Temperature	Humidity	Moisture	Nitrogen	Potassium	Phosphorous
count	99.000000	99.000000	99.000000	99.000000	99.000000	99.000000
mean	30.282828	59.151515	43.181818	18.909091	3.383838	18.606061
std	3.502304	5.840331	11.271568	11.599693	5.814667	13.476978
min	25.000000	50.000000	25.000000	4.000000	0.000000	0.000000
25%	28.000000	54.000000	34.000000	10.000000	0.000000	9.000000
50%	30.000000	60.000000	41.000000	13.000000	0.000000	19.000000
75%	33.000000	64.000000	50.500000	24.000000	7.500000	30.000000
max	38.000000	72.000000	65.000000	42.000000	19.000000	42.000000

Descriptive statistics (categorical features):


	Soil Type	Crop Type	Fertilizer Name
count	99	99	99
unique	5	11	7
top	Loamy	Sugarcane	Urea
freq	21	13	22


Fertilizer Name Distribution:

Fertilizer Name	count
Urea	22
DAP	18
28-28	17
14-35-14	14
20-20	14
17-17-17	7
10-26-26	7

Name: count, dtype: int64

Distribution of Fertilizer Names





Soil Type Distribution:

Soil Type

Loamy 21

Sandy 20


Clayey 20

Black 19

Red 19

Name: count, dtype: int64

Distribution of Soil Types



Crop Type Distribution:

Crop Type

Sugarcane 13

Cotton 12

Milletts 11

Pulses 10

Paddy 10

Wheat 9

Barley 7

Oil seeds 7


Tobacco 7

Ground Nuts 7


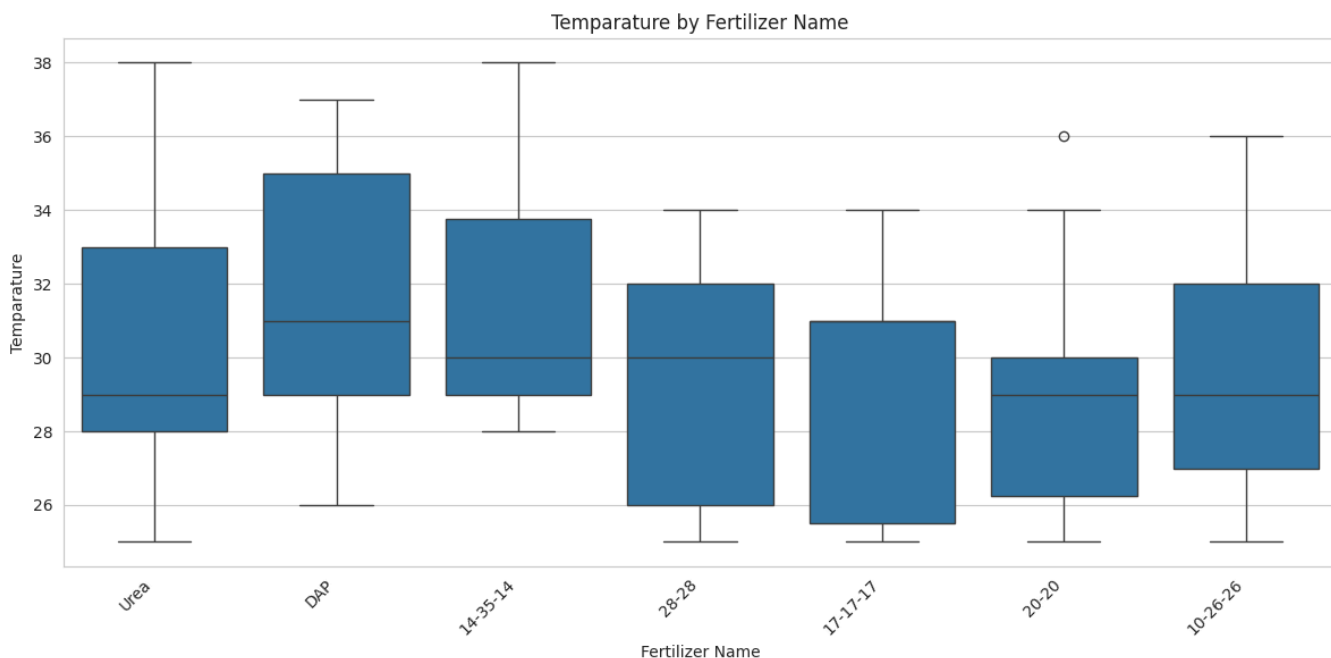
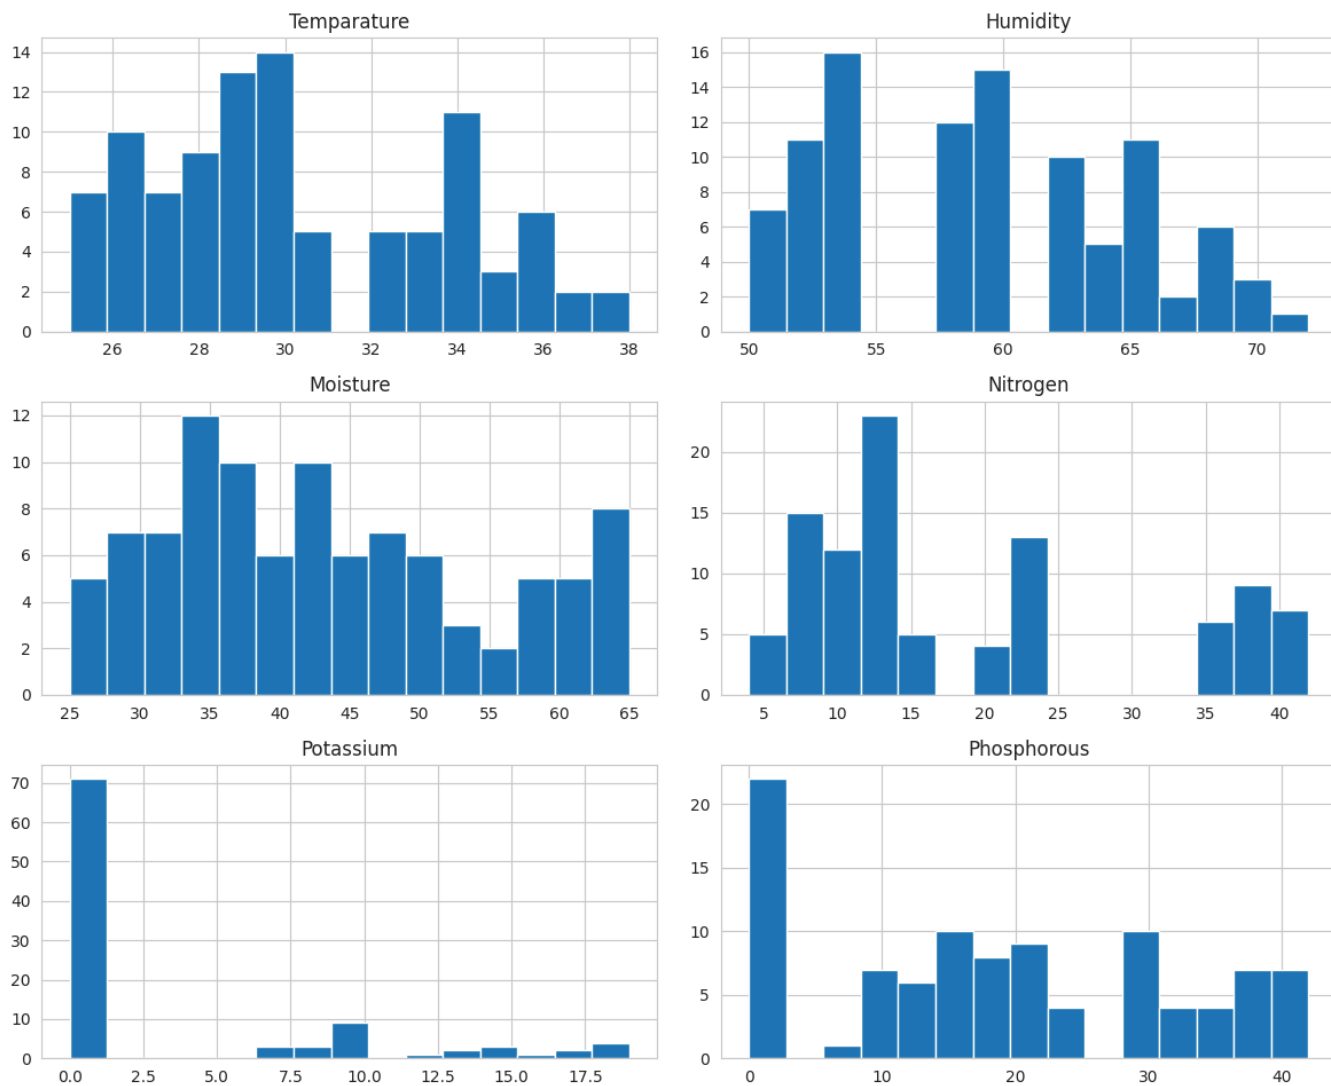
Maize 6


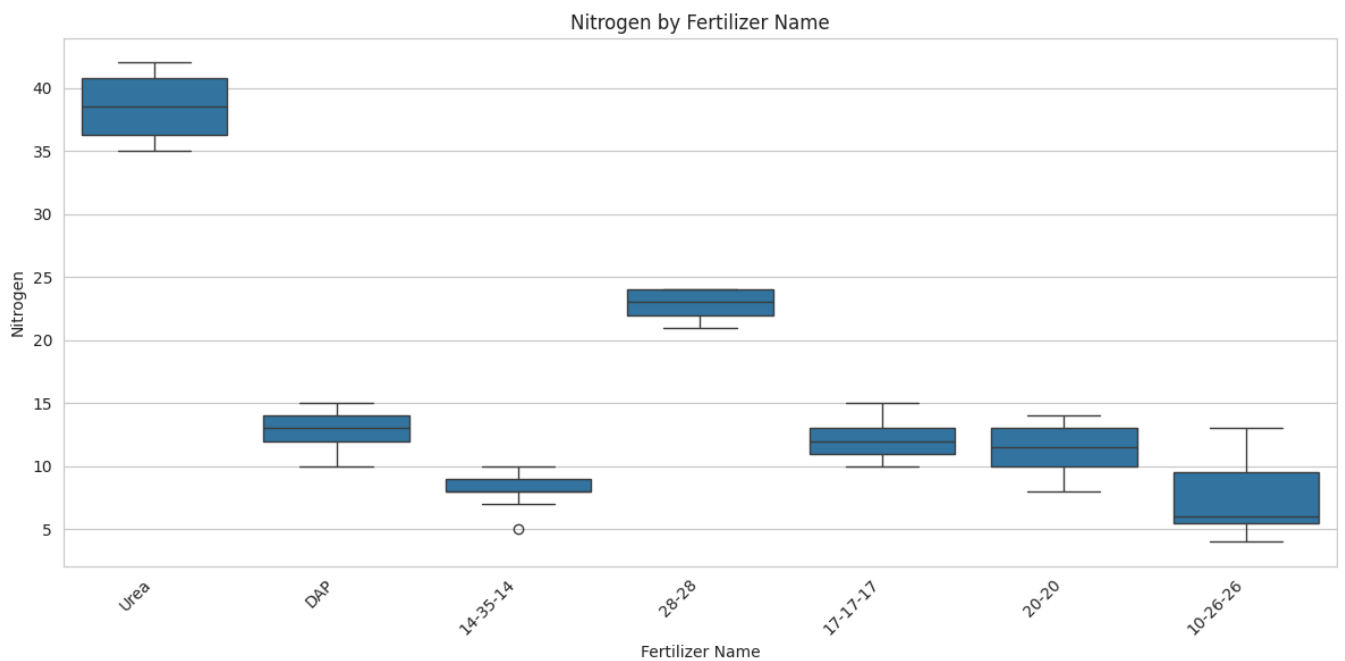
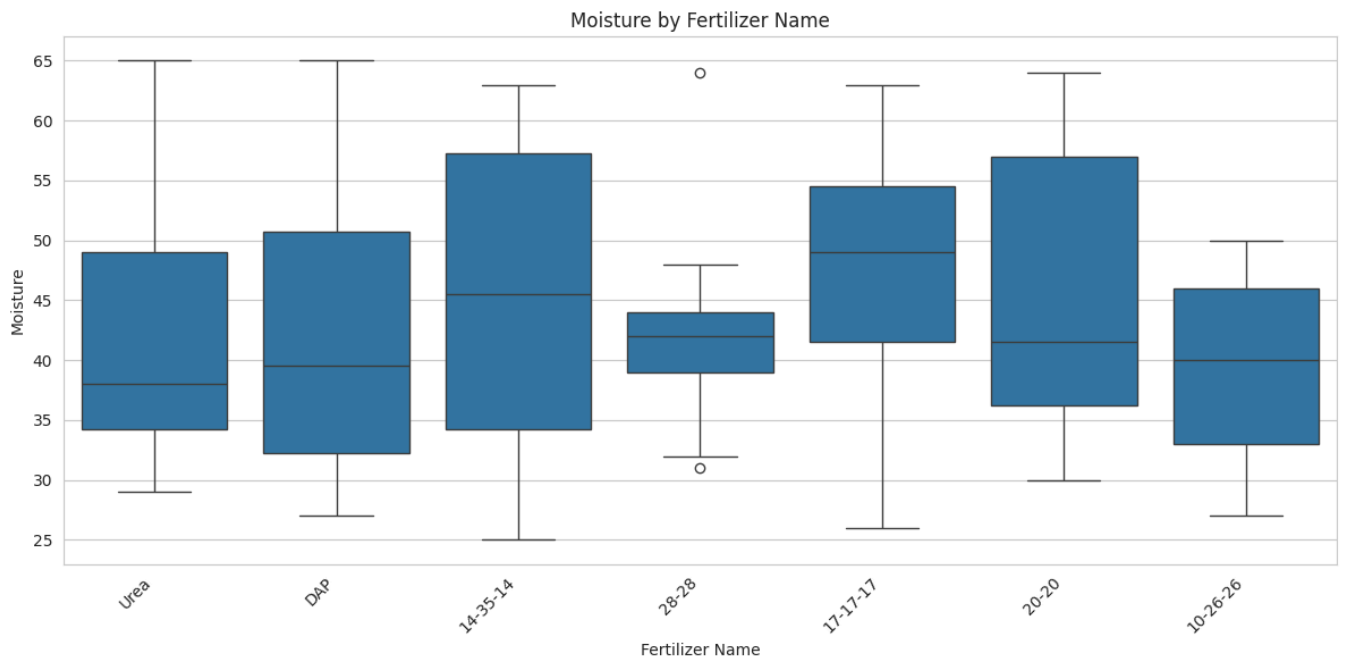

Name: count, dtype: int64


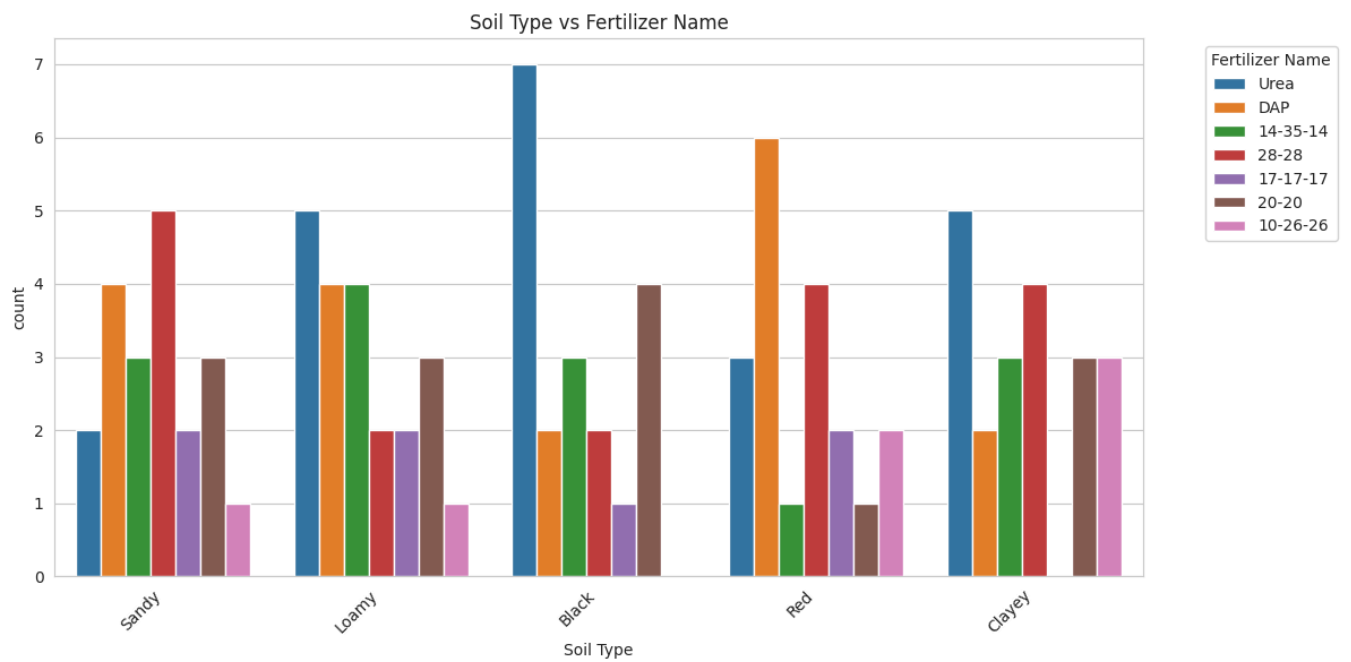
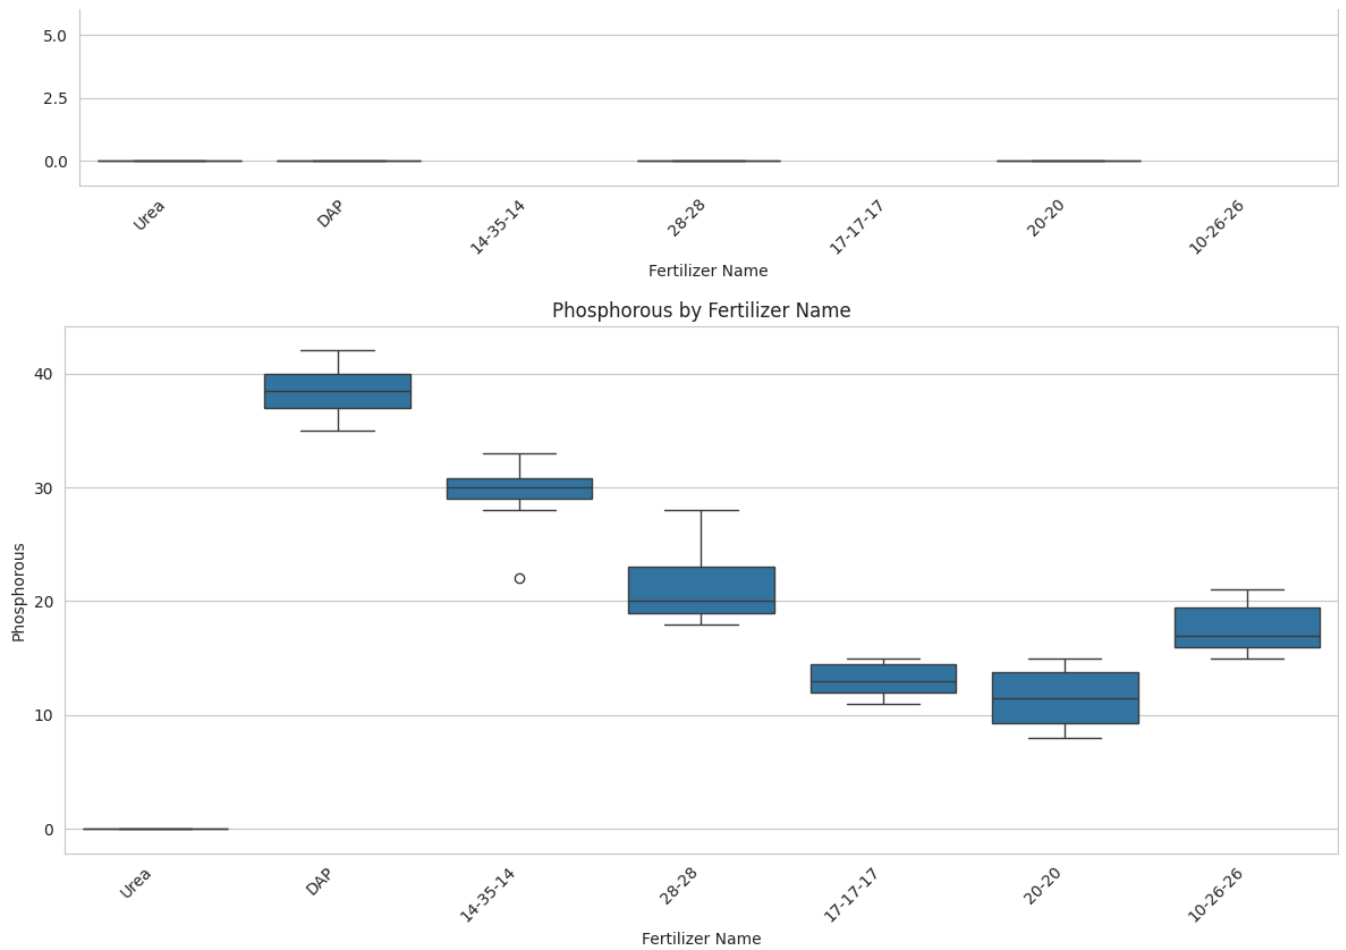
Distribution of Crop Types (Fertilizer Data)




Histograms of Numerical Features (Fertilizer Data)









2. Data Preprocessing for Fertilizer Recommendation

```

if 'fertilizer_df' in locals():
    X_fert = fertilizer_df.drop('Fertilizer Name', axis=1)
    y_fert_str = fertilizer_df['Fertilizer Name']

    # Encode the target variable (Fertilizer Name)
    label_encoder_fert = LabelEncoder()
    y_fert = label_encoder_fert.fit_transform(y_fert_str)
    print("\nEncoded Fertilizer Labels Mapping:")
    for i, class_name in enumerate(label_encoder_fert.classes_):
        print(f"{class_name}: {i}")

    # Identify numerical and categorical columns
    numerical_features_fert = X_fert.select_dtypes(include=np.number).columns.tolist()
    categorical_features_fert = X_fert.select_dtypes(include='object').columns.tolist()

    print("\nNumerical features for fertilizer:", numerical_features_fert)
    print("Categorical features for fertilizer:", categorical_features_fert)

    # Create preprocessor
    # Numerical pipeline: Standard scaling
    # Categorical pipeline: One-hot encoding
    preprocessor_fert = ColumnTransformer(
        transformers=[
            ('num', StandardScaler(), numerical_features_fert),
            ('cat', OneHotEncoder(handle_unknown='ignore', drop='first'), categorical_features_fert)
        ],
        remainder='passthrough' # Keep any other columns (should be none here)
    )

    # Split data
    X_train_fert, X_test_fert, y_train_fert, y_test_fert = train_test_split(
        X_fert, y_fert, test_size=0.2, random_state=42, stratify=y_fert
    )

    # Apply preprocessing
    # Fit on training data and transform both train and test
    X_train_fert_processed = preprocessor_fert.fit_transform(X_train_fert)
    X_test_fert_processed = preprocessor_fert.transform(X_test_fert)

    # Get feature names after one-hot encoding for interpretability later (optional)
    try:
        ohe_feature_names = preprocessor_fert.named_transformers_['cat'].get_feature_names_out(categorical_features_fert)
        all_feature_names_fert = numerical_features_fert + list(ohe_feature_names)
        print("\nTotal features after processing:", len(all_feature_names_fert))
    except Exception as e:
        print(f"Could not get OHE feature names: {e}")

    print("\nShape of X_train_fert_processed:", X_train_fert_processed.shape)
    print("Shape of X_test_fert_processed:", X_test_fert_processed.shape)
else:
    print("fertilizer_df not loaded. Skipping Preprocessing for Fertilizer Recommendation.")

```



Encoded Fertilizer Labels Mapping:

```

10-26-26: 0
14-35-14: 1
17-17-17: 2
20-20: 3
28-28: 4
DAP: 5
Urea: 6

```

```

Numerical features for fertilizer: ['Temperature', 'Humidity', 'Moisture', 'Nitrogen', 'Potassium', 'Phosphorous']
Categorical features for fertilizer: ['Soil Type', 'Crop Type']

```

Total features after processing: 20

```

Shape of X_train_fert_processed: (79, 20)
Shape of X_test_fert_processed: (20, 20)

```

3. Model Selection & Training for Fertilizer Recommendation

```

if 'fertilizer_df' in locals():
    # Using similar models as for crop recommendation
    models_fert = {
        "Logistic Regression": LogisticRegression(max_iter=1000, solver='liblinear', multi_class='ovr', random_state=42),
        "K-Nearest Neighbors": KNeighborsClassifier(),
    }

```

```

    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "Random Forest": RandomForestClassifier(random_state=42),
    "Gradient Boosting": GradientBoostingClassifier(random_state=42)
}

results_fert = {}
print("\n--- Training Fertilizer Recommendation Models ---")

for name, model in models_fert.items():
    # Create a full pipeline for each model to ensure preprocessing is part of CV if used later
    # For now, we use pre-processed data directly
    model.fit(X_train_fert_processed, y_train_fert)
    y_pred_fert = model.predict(X_test_fert_processed)

    accuracy = accuracy_score(y_test_fert, y_pred_fert)
    f1 = f1_score(y_test_fert, y_pred_fert, average='weighted')
    results_fert[name] = {'Accuracy': accuracy, 'F1-score (Weighted)': f1}

    print(f"\n--- {name} (Fertilizer Recommendation) ---")
    print(f"Accuracy: {accuracy:.4f}")
    print(f"F1-score (Weighted): {f1:.4f}")

fert_results_df = pd.DataFrame(results_fert).T.sort_values(by='Accuracy', ascending=False)
print("\n--- Model Performance Summary (Fertilizer Recommendation) ---")
print(fert_results_df)
else:
    print("fertilizer_df not loaded. Skipping Model Training for Fertilizer Recommendation.")

```

```

--- Training Fertilizer Recommendation Models ---

--- Logistic Regression (Fertilizer Recommendation) ---
Accuracy: 0.9500
F1-score (Weighted): 0.9333

--- K-Nearest Neighbors (Fertilizer Recommendation) ---
Accuracy: 0.9000
F1-score (Weighted): 0.8986

--- Decision Tree (Fertilizer Recommendation) ---
Accuracy: 0.9500
F1-score (Weighted): 0.9533
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:1256: FutureWarning: 'multi_class' was deprecated in versi
warnings.warn(

--- Random Forest (Fertilizer Recommendation) ---
Accuracy: 1.0000
F1-score (Weighted): 1.0000

--- Gradient Boosting (Fertilizer Recommendation) ---
Accuracy: 0.9500
F1-score (Weighted): 0.9473

--- Model Performance Summary (Fertilizer Recommendation) ---
              Accuracy  F1-score (Weighted)
Random Forest          1.00          1.000000
Logistic Regression     0.95          0.933333
Decision Tree           0.95          0.953333
Gradient Boosting       0.95          0.947273
K-Nearest Neighbors     0.90          0.898571

```

4. Evaluation for Fertilizer Recommendation

```

if 'fertilizer_df' in locals():
    best_model_name_fert = fert_results_df.index[0]
    best_model_fert = models_fert[best_model_name_fert] # The trained model instance


    y_pred_best_fert = best_model_fert.predict(X_test_fert_processed)

    print(f"\n--- Detailed Evaluation for {best_model_name_fert} (Fertilizer Recommendation) ---")
    print("Classification Report:\n", classification_report(y_test_fert, y_pred_best_fert, target_names=label_encoder_fert.classes_, zero

    cm_fert = confusion_matrix(y_test_fert, y_pred_best_fert)
    plt.figure(figsize=(10, 8))
    sns.heatmap(cm_fert, annot=True, fmt='d', cmap='Greens',
                xticklabels=label_encoder_fert.classes_,
                yticklabels=label_encoder_fert.classes_)
    plt.title(f'Confusion Matrix for {best_model_name_fert} (Fertilizer Recommendation)')
    plt.xlabel('Predicted Label')
    plt.ylabel('True Label')
    plt.xticks(rotation=45, ha='right')
    plt.yticks(rotation=0)
    plt.tight_layout()

```


```
plt.show()
else:
    print("fertilizer_df not loaded. Skipping Evaluation for Fertilizer Recommendation.")
```



--- Detailed Evaluation for Random Forest (Fertilizer Recommendation) ---

Classification Report:

	precision	recall	f1-score	support
10-26-26	1.00	1.00	1.00	1
14-35-14	1.00	1.00	1.00	3
17-17-17	1.00	1.00	1.00	1
20-20	1.00	1.00	1.00	3
28-28	1.00	1.00	1.00	3
DAP	1.00	1.00	1.00	4
Urea	1.00	1.00	1.00	5
accuracy			1.00	20
macro avg	1.00	1.00	1.00	20
weighted avg	1.00	1.00	1.00	20



✓ THANK YOU

Start coding or [generate](#) with AI.

Double-click (or enter) to edit